

weatheripynb

February 27, 2024

```
[1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import LabelEncoder
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression

from sklearn.metrics import r2_score
from sklearn.preprocessing import PowerTransformer
from sklearn.model_selection import cross_val_score
from scipy.cluster.vq import whiten, kmeans, vq
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score

import warnings
warnings.filterwarnings('ignore')
```

```
[2]: df = pd.read_csv('weather.csv')
```

```
[3]: df
```

```
[3]:
```

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	\
0	8.0	24.3	0.0	3.4	6.3	NW	
1	14.0	26.9	3.6	4.4	9.7	ENE	
2	13.7	23.4	3.6	5.8	3.3	NW	
3	13.3	15.5	39.8	7.2	9.1	NW	
4	7.6	16.1	2.8	5.6	10.6	SSE	
..		
361	9.0	30.7	0.0	7.6	12.1	NNW	
362	7.1	28.4	0.0	11.6	12.7	N	
363	12.5	19.9	0.0	8.4	5.3	ESE	
364	12.5	26.9	0.0	5.0	7.1	NW	
365	12.3	30.2	0.0	6.0	12.6	NW	

WindGustSpeed WindDir9am WindDir3pm WindSpeed9am ... Humidity3pm \

0	30.0	SW	NW	6.0	...	29
1	39.0	E	W	4.0	...	36
2	85.0	N	NNE	6.0	...	69
3	54.0	WNW	W	30.0	...	56
4	50.0	SSE	ESE	20.0	...	49
..
361	76.0	SSE	NW	7.0	...	15
362	48.0	NNW	NNW	2.0	...	22
363	43.0	ENE	ENE	11.0	...	47
364	46.0	SSW	WNW	6.0	...	39
365	78.0	NW	WNW	31.0	...	13

	Pressure9am	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm	\
0	1019.7	1015.0	7	7	14.4	23.6	
1	1012.4	1008.4	5	3	17.5	25.7	
2	1009.5	1007.2	8	7	15.4	20.2	
3	1005.5	1007.0	2	7	13.5	14.1	
4	1018.3	1018.5	7	7	11.1	15.4	
..	
361	1016.1	1010.8	1	3	20.4	30.0	
362	1020.0	1016.9	0	1	17.2	28.2	
363	1024.0	1022.8	3	2	14.5	18.3	
364	1021.0	1016.2	6	7	15.8	25.9	
365	1009.6	1009.2	1	1	23.8	28.6	

	RainToday	RISK_MM	RainTomorrow
0	No	3.6	Yes
1	Yes	3.6	Yes
2	Yes	39.8	Yes
3	Yes	2.8	Yes
4	Yes	0.0	No
..
361	No	0.0	No
362	No	0.0	No
363	No	0.0	No
364	No	0.0	No
365	No	0.0	No

[366 rows x 22 columns]

```
[4]: obj_df = df.select_dtypes(include=['object'])
      num_df = df.select_dtypes([np.number])
```

```
[5]: obj_df.fillna(obj_df.mode().iloc[0], inplace=True)
```

```
[6]: obj_df_label = obj_df.copy()
```

```
[7]: le = LabelEncoder()
obj_df_label = obj_df.apply(le.fit_transform)
```

```
[8]: obj_df_label
```

```
[8]:      WindGustDir  WindDir9am  WindDir3pm  RainToday  RainTomorrow
0              7           12           7           0           1
1              1            0          13           1           1
2              7            3            5           1           1
3              7           14          13           1           1
4             10           10            2           1           0
..          ...          ...          ...          ...          ...
361             6           10            7           0           0
362             3            6            6           0           0
363             2            1            1           0           0
364             7           11          14           0           0
365             7            7          14           0           0
```

[366 rows x 5 columns]

```
[9]: obj_df_onehot = obj_df.copy()
```

```
[10]: obj_df_onehot = pd.get_dummies(obj_df_onehot, columns=obj_df_onehot.columns)
```

```
[11]: obj_df_onehot
```

```
[11]:      WindGustDir_E  WindGustDir_ENE  WindGustDir_ESE  WindGustDir_N  \
0              False              False              False              False
1              False               True              False              False
2              False              False              False              False
3              False              False              False              False
4              False              False              False              False
..          ...          ...          ...          ...
361             False              False              False              False
362             False              False              False              True
363             False              False              True              False
364             False              False              False              False
365             False              False              False              False

      WindGustDir_NE  WindGustDir_NNE  WindGustDir_NNW  WindGustDir_NW  \
0              False              False              False              True
1              False              False              False              False
2              False              False              False              True
3              False              False              False              True
4              False              False              False              False
..          ...          ...          ...          ...
361             False              False              True              False
```

362	False	False	False	False
363	False	False	False	False
364	False	False	False	True
365	False	False	False	True

	WindGustDir_S	WindGustDir_SE	...	WindDir3pm_SSE	WindDir3pm_SSW	\
0	False	False	...	False	False	
1	False	False	...	False	False	
2	False	False	...	False	False	
3	False	False	...	False	False	
4	False	False	...	False	False	
..	
361	False	False	...	False	False	
362	False	False	...	False	False	
363	False	False	...	False	False	
364	False	False	...	False	False	
365	False	False	...	False	False	

	WindDir3pm_SW	WindDir3pm_W	WindDir3pm_WNW	WindDir3pm_WSW	\
0	False	False	False	False	
1	False	True	False	False	
2	False	False	False	False	
3	False	True	False	False	
4	False	False	False	False	
..	
361	False	False	False	False	
362	False	False	False	False	
363	False	False	False	False	
364	False	False	True	False	
365	False	False	True	False	

	RainToday_No	RainToday_Yes	RainTomorrow_No	RainTomorrow_Yes
0	True	False	False	True
1	False	True	False	True
2	False	True	False	True
3	False	True	False	True
4	False	True	True	False
..
361	True	False	True	False
362	True	False	True	False
363	True	False	True	False
364	True	False	True	False
365	True	False	True	False

[366 rows x 52 columns]

```
[12]: obj_df_binary = obj_df.copy()
```

```
[13]: obj_df.isnull().sum()
```

```
[13]: WindGustDir      0
      WindDir9am      0
      WindDir3pm      0
      RainToday       0
      RainTomorrow    0
      dtype: int64
```

```
[14]: obj_df_binary.fillna(obj_df_binary.mode().iloc[0], inplace=True)
```

```
[15]: obj_df_binary
```

```
[15]:
```

	WindGustDir	WindDir9am	WindDir3pm	RainToday	RainTomorrow
0	NW	SW	NW	No	Yes
1	ENE	E	W	Yes	Yes
2	NW	N	NNE	Yes	Yes
3	NW	WNW	W	Yes	Yes
4	SSE	SSE	ESE	Yes	No
..
361	NNW	SSE	NW	No	No
362	N	NNW	NNW	No	No
363	ESE	ENE	ENE	No	No
364	NW	SSW	WNW	No	No
365	NW	NW	WNW	No	No

```
[366 rows x 5 columns]
```

```
[16]: # lb = LabelBinarizer()
      # obj_df_binary = lb.fit_transform(obj_df_binary)
```

```
[17]: num_df.isnull().sum()
```

```
[17]: MinTemp      0
      MaxTemp      0
      Rainfall    0
      Evaporation  0
      Sunshine    3
      WindGustSpeed 2
      WindSpeed9am 7
      WindSpeed3pm 0
      Humidity9am   0
      Humidity3pm   0
      Pressure9am   0
      Pressure3pm   0
      Cloud9am      0
      Cloud3pm      0
```

```
Temp9am      0
Temp3pm      0
RISK_MM      0
dtype: int64
```

```
[18]: num_df
```

```
[18]:
```

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	\
0	8.0	24.3	0.0	3.4	6.3	30.0	
1	14.0	26.9	3.6	4.4	9.7	39.0	
2	13.7	23.4	3.6	5.8	3.3	85.0	
3	13.3	15.5	39.8	7.2	9.1	54.0	
4	7.6	16.1	2.8	5.6	10.6	50.0	
..	
361	9.0	30.7	0.0	7.6	12.1	76.0	
362	7.1	28.4	0.0	11.6	12.7	48.0	
363	12.5	19.9	0.0	8.4	5.3	43.0	
364	12.5	26.9	0.0	5.0	7.1	46.0	
365	12.3	30.2	0.0	6.0	12.6	78.0	

	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity3pm	Pressure9am	\
0	6.0	20	68	29	1019.7	
1	4.0	17	80	36	1012.4	
2	6.0	6	82	69	1009.5	
3	30.0	24	62	56	1005.5	
4	20.0	28	68	49	1018.3	
..	
361	7.0	50	38	15	1016.1	
362	2.0	19	45	22	1020.0	
363	11.0	9	63	47	1024.0	
364	6.0	28	69	39	1021.0	
365	31.0	35	43	13	1009.6	

	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RISK_MM
0	1015.0	7	7	14.4	23.6	3.6
1	1008.4	5	3	17.5	25.7	3.6
2	1007.2	8	7	15.4	20.2	39.8
3	1007.0	2	7	13.5	14.1	2.8
4	1018.5	7	7	11.1	15.4	0.0
..
361	1010.8	1	3	20.4	30.0	0.0
362	1016.9	0	1	17.2	28.2	0.0
363	1022.8	3	2	14.5	18.3	0.0
364	1016.2	6	7	15.8	25.9	0.0
365	1009.2	1	1	23.8	28.6	0.0

```
[366 rows x 17 columns]
```

```
[19]: num_df_mean = num_df.copy()
```

```
[20]: num_df_mean.fillna(num_df_mean.mean().iloc[0], inplace=True)
```

```
[21]: num_df_mean.isnull().sum()
```

```
[21]: MinTemp      0
      MaxTemp      0
      Rainfall     0
      Evaporation  0
      Sunshine     0
      WindGustSpeed 0
      WindSpeed9am 0
      WindSpeed3pm 0
      Humidity9am   0
      Humidity3pm   0
      Pressure9am   0
      Pressure3pm   0
      Cloud9am      0
      Cloud3pm      0
      Temp9am       0
      Temp3pm       0
      RISK_MM       0
      dtype: int64
```

```
[22]: obj_df_label
```

```
[22]:      WindGustDir  WindDir9am  WindDir3pm  RainToday  RainTomorrow
0              7           12           7           0             1
1              1            0          13           1             1
2              7            3            5           1             1
3              7           14          13           1             1
4             10           10            2           1             0
..           ...           ...           ...           ...           ...
361             6           10            7           0             0
362             3            6            6           0             0
363             2            1            1           0             0
364             7           11          14           0             0
365             7            7          14           0             0
```

[366 rows x 5 columns]

```
[23]: # Q1 = num_df_mean.quantile(0.25)
      # Q3 = num_df_mean.quantile(0.75)
      # IQR = Q3 - Q1
      # lower_fence = Q1 - 1.5 * IQR
      # upper_fence = Q3 + 1.5 * IQR
```

```
# num_df_mean = num_df_mean[~((num_df_mean < lower_fence) | (num_df_mean > upper_fence)).any(axis=1)]
```

```
[24]: processed_df = pd.concat([num_df_mean, obj_df_label], axis=1)
```

```
[25]: processed_df
```

```
[25]:
```

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	\
0	8.0	24.3	0.0	3.4	6.3	30.0	
1	14.0	26.9	3.6	4.4	9.7	39.0	
2	13.7	23.4	3.6	5.8	3.3	85.0	
3	13.3	15.5	39.8	7.2	9.1	54.0	
4	7.6	16.1	2.8	5.6	10.6	50.0	
..	
361	9.0	30.7	0.0	7.6	12.1	76.0	
362	7.1	28.4	0.0	11.6	12.7	48.0	
363	12.5	19.9	0.0	8.4	5.3	43.0	
364	12.5	26.9	0.0	5.0	7.1	46.0	
365	12.3	30.2	0.0	6.0	12.6	78.0	

	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity3pm	...	Cloud9am	\
0	6.0	20	68	29	...	7	
1	4.0	17	80	36	...	5	
2	6.0	6	82	69	...	8	
3	30.0	24	62	56	...	2	
4	20.0	28	68	49	...	7	
..	
361	7.0	50	38	15	...	1	
362	2.0	19	45	22	...	0	
363	11.0	9	63	47	...	3	
364	6.0	28	69	39	...	6	
365	31.0	35	43	13	...	1	

	Cloud3pm	Temp9am	Temp3pm	RISK_MM	WindGustDir	WindDir9am	WindDir3pm	\
0	7	14.4	23.6	3.6	7	12	7	
1	3	17.5	25.7	3.6	1	0	13	
2	7	15.4	20.2	39.8	7	3	5	
3	7	13.5	14.1	2.8	7	14	13	
4	7	11.1	15.4	0.0	10	10	2	
..	
361	3	20.4	30.0	0.0	6	10	7	
362	1	17.2	28.2	0.0	3	6	6	
363	2	14.5	18.3	0.0	2	1	1	
364	7	15.8	25.9	0.0	7	11	14	
365	1	23.8	28.6	0.0	7	7	14	

```
RainToday RainTomorrow
```



```

0          0          1
1          1          1
2          1          1
3          1          1
4          1          0
..         ...         ...
361        0          0
362        0          0
363        0          0
364        0          0
365        0          0

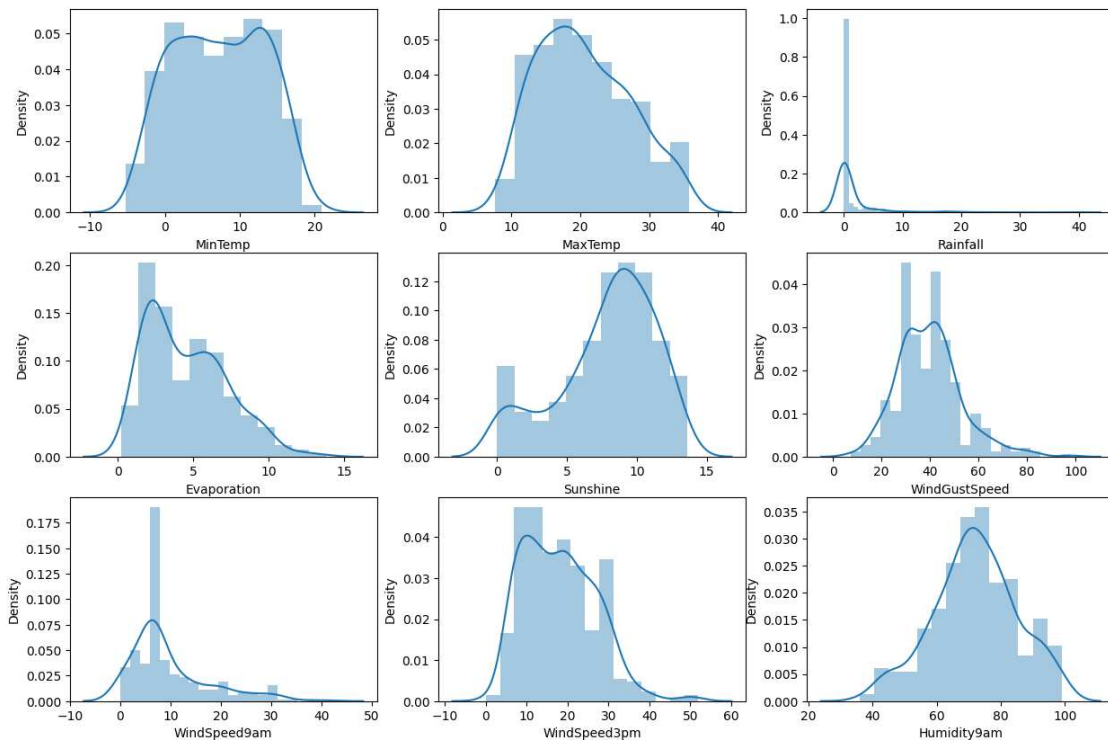
```

[366 rows x 22 columns]

```

[81]: fig, ax = plt.subplots(3, 3, figsize=(15, 10))
      for i, subplot in zip(processed_df.columns, ax.flatten()):
          sns.distplot(processed_df[i], ax=subplot)
      plt.show()

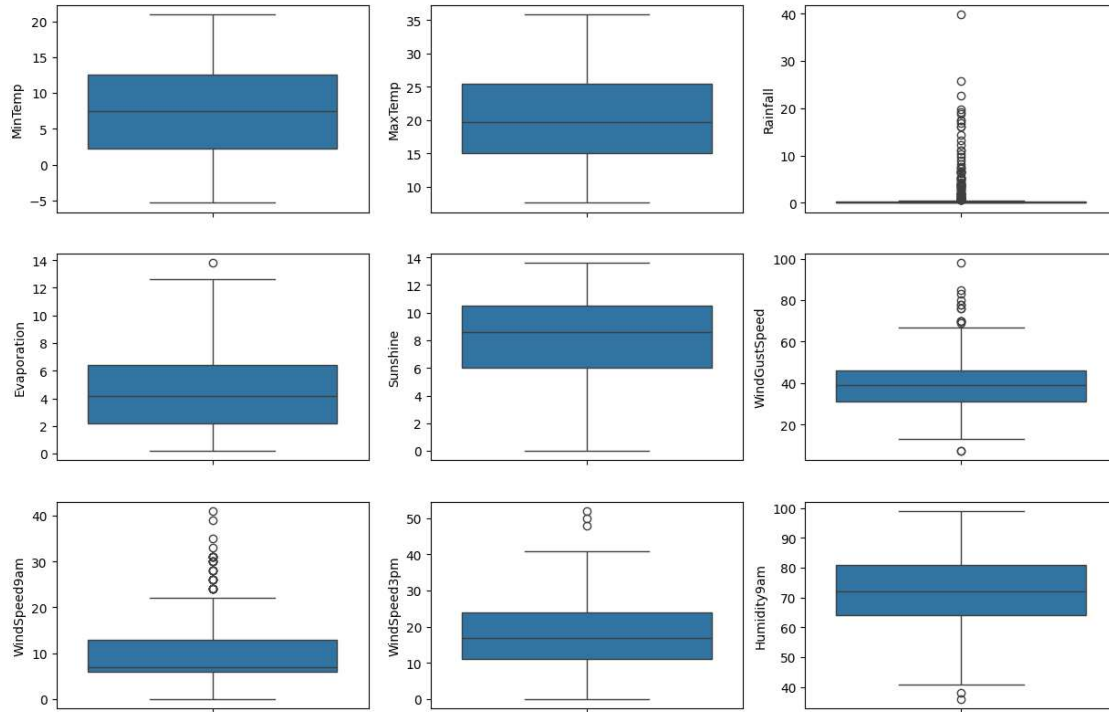
```



```

[80]: fig, ax = plt.subplots(3, 3, figsize=(15, 10))
      for i, subplot in zip(processed_df.columns, ax.flatten()):
          sns.boxplot(processed_df[i], ax=subplot)
      plt.show()


```



```
[28]: processed_df_log = processed_df.copy()
```

```
[29]: processed_df_log = np.log(processed_df_log)
processed_df_log = processed_df_log.replace([np.inf, -np.inf], np.nan)
np.seterr(divide = 'ignore')
warnings.filterwarnings('ignore')
```


```
[30]: sns.distplot(processed_df_log)
```



```
[31]: processed_df_sqrt = processed_df.copy()
```

```
[32]: processed_df_sqrt = np.sqrt(processed_df_sqrt)
```

```
[33]: sns.distplot(processed_df_sqrt)  
warnings.filterwarnings('ignore')
```




```
[34]: processed_df_reciprocal = processed_df.copy()
```

```
[35]: processed_df_reciprocal = 1/processed_df_reciprocal  
processed_df_reciprocal = processed_df_reciprocal.replace([np.inf, -np.inf], np.  
↪ nan)
```

```
[36]: sns.distplot(processed_df_reciprocal)
```

```
[36]: <Axes: ylabel='Density'>
```



```
[86]: processed_df.head()
```

```
[86]:
```

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	\
0	8.0	24.3	0.0	3.4	6.3	30.0	
1	14.0	26.9	3.6	4.4	9.7	39.0	
2	13.7	23.4	3.6	5.8	3.3	85.0	
3	13.3	15.5	39.8	7.2	9.1	54.0	
4	7.6	16.1	2.8	5.6	10.6	50.0	

	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity3pm	...	Cloud9am	\
0	6.0	20	68	29	...	7	
1	4.0	17	80	36	...	5	
2	6.0	6	82	69	...	8	
3	30.0	24	62	56	...	2	
4	20.0	28	68	49	...	7	

	Cloud3pm	Temp9am	Temp3pm	RISK_MM	WindGustDir	WindDir9am	WindDir3pm	\
0	7	14.4	23.6	3.6	7	12	7	
1	3	17.5	25.7	3.6	1	0	13	
2	7	15.4	20.2	39.8	7	3	5	
3	7	13.5	14.1	2.8	7	14	13	
4	7	11.1	15.4	0.0	10	10	2	

	RainToday	RainTomorrow
0	0	1
1	1	1
2	1	1
3	1	1
4	1	0

[5 rows x 22 columns]

```
[38]: lr = LinearRegression()
```

```
[39]: X = processed_df.iloc[:, :-5]
      y = processed_df.iloc[:, -5:]
```

```
[40]: y
```

```
[40]:
```

	WindGustDir	WindDir9am	WindDir3pm	RainToday	RainTomorrow
0	7	12	7	0	1
1	1	0	13	1	1
2	7	3	5	1	1
3	7	14	13	1	1
4	10	10	2	1	0
...
361	6	10	7	0	0
362	3	6	6	0	0
363	2	1	1	0	0
364	7	11	14	0	0
365	7	7	14	0	0

[366 rows x 5 columns]

```
[41]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2)
```

```
[42]: lr.fit(X_train,y_train)

      y_pred = lr.predict(X_test)

      r2_score(y_test,y_pred)
```

```
[42]: 0.27407628542111473
```

```
[43]: pt1 = PowerTransformer()

      X_train_transformed2 = pt1.fit_transform(X_train)
      X_test_transformed2 = pt1.transform(X_test)
```

```

lr = LinearRegression()
lr.fit(X_train_transformed2,y_train)

y_pred3 = lr.predict(X_test_transformed2)

print(r2_score(y_test,y_pred3))

pd.DataFrame({'cols':X_train.columns,'Yeo_Johnson_lambdas':pt1.lambdas_})

```

0.36108092627846095

```

[43]:
      cols  Yeo_Johnson_lambdas
0    MinTemp      0.845582
1    MaxTemp      0.258255
2    Rainfall     -2.494539
3  Evaporation      0.146360
4    Sunshine      1.432938
5  WindGustSpeed      0.519691
6  WindSpeed9am      0.255951
7  WindSpeed3pm      0.452736
8    Humidity9am      1.266156
9    Humidity3pm      0.372768
10  Pressure9am     25.954197
11  Pressure3pm     22.063945
12    Cloud9am      0.299128
13    Cloud3pm      0.394528
14    Temp9am       0.811665
15    Temp3pm       0.426617
16    RISK_MM      -2.305938

```

```

[44]: pt = PowerTransformer()
      X_transformed2 = pt.fit_transform(X)

      lr = LinearRegression()
      np.mean(cross_val_score(lr,X_transformed2,y,scoring='r2'))

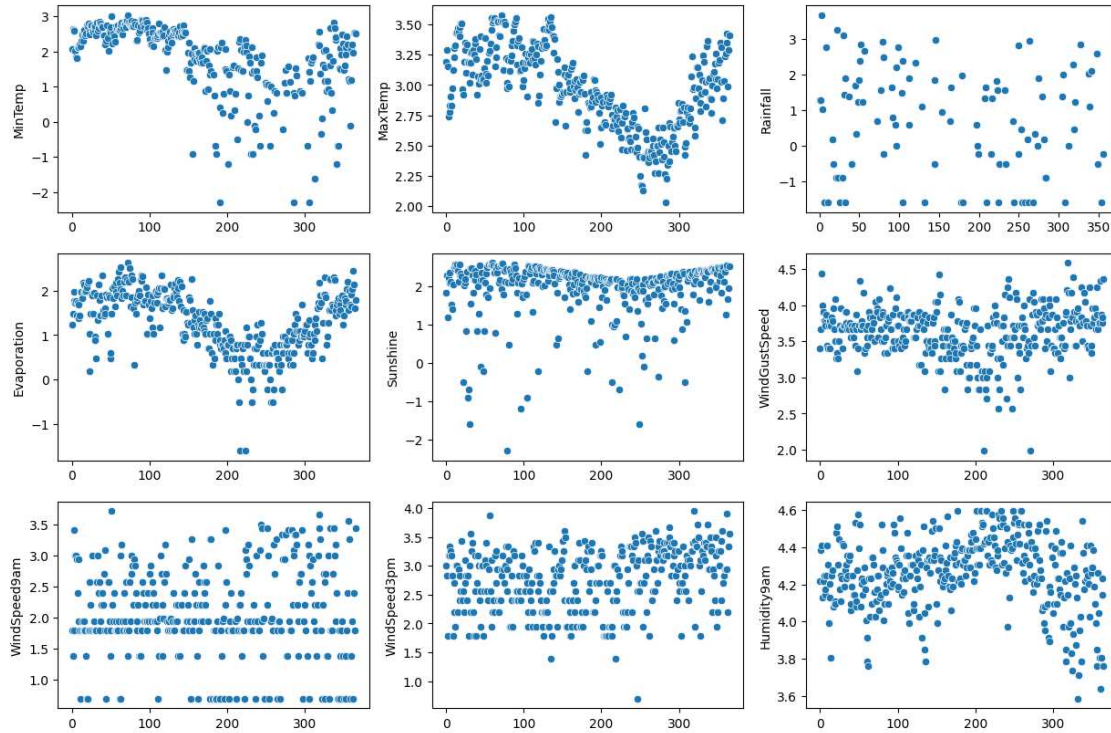
```

[44]: 0.30544757696266495

```


[79]: fig, ax = plt.subplots(3, 3, figsize=(15, 10))
      for i, subplot in zip(processed_df_log.columns, ax.flatten()):
          sns.scatterplot(processed_df_log[i], ax=subplot)
      plt.show()

```

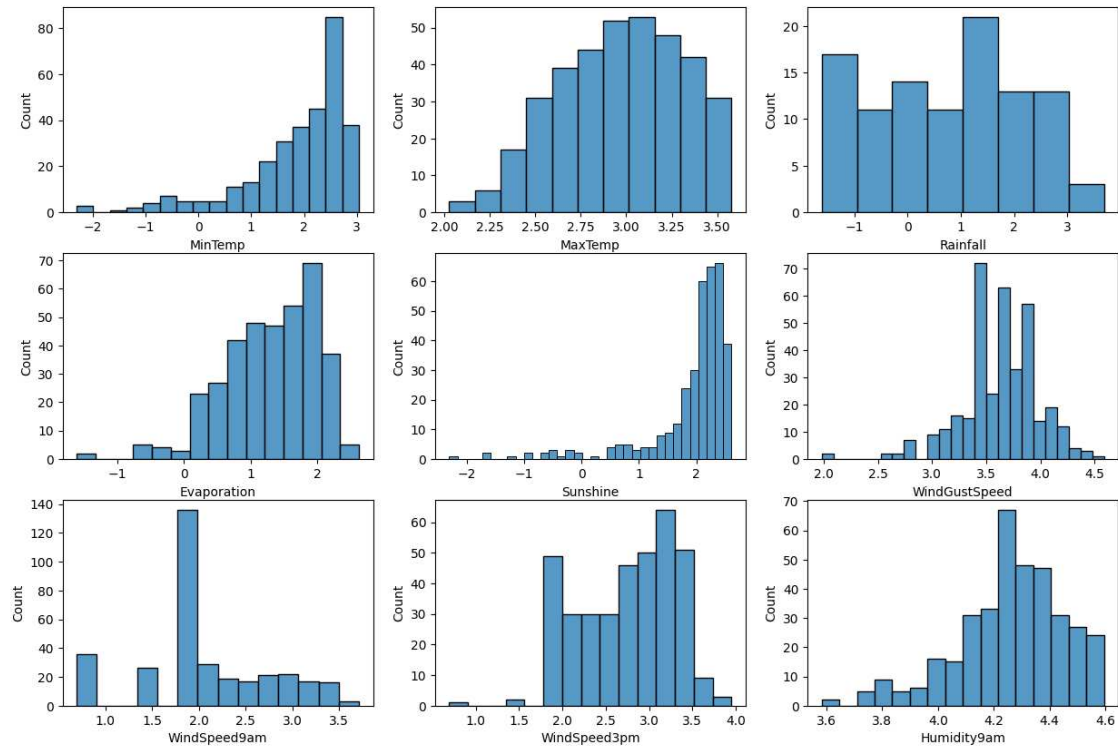


```
[46]: sns.heatmap(processed_df_log.corr(), annot=True)
```

```
[46]: <Axes: >
```

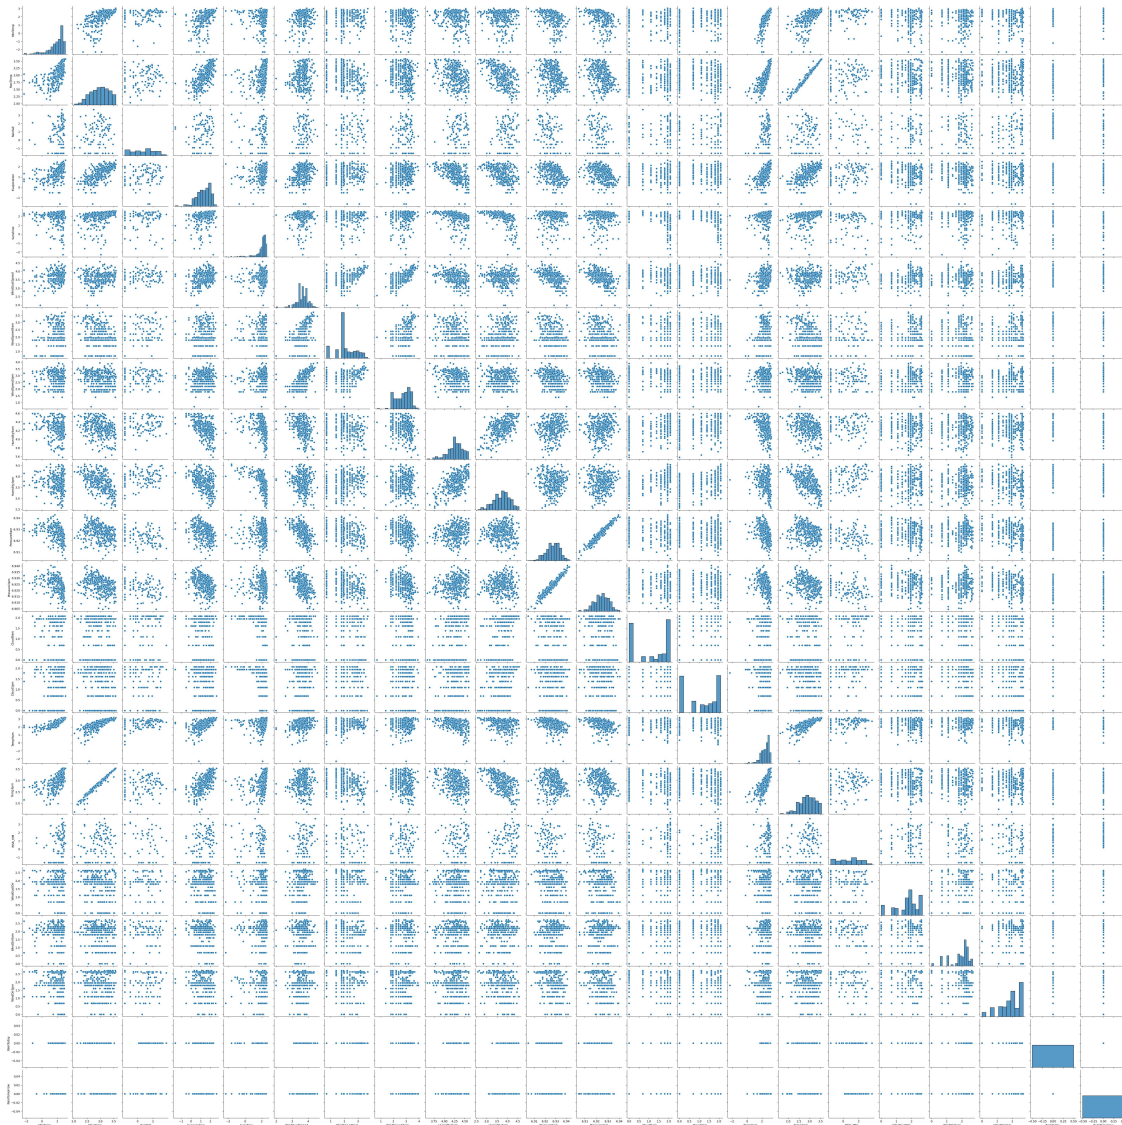



```
[72]: fig, ax = plt.subplots(3, 3, figsize=(15, 10))
      for i, subplot in zip(processed_df_log.columns, ax.flatten()):
          sns.histplot(processed_df_log[i], ax=subplot)
      plt.show()
```



```
[48]: sns.pairplot(processed_df_log)
```

```
[48]: <seaborn.axisgrid.PairGrid at 0x708591164b10>
```



```
[49]: processed_df_log.columns
```

```
[49]: Index(['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'Sunshine',
          'WindGustSpeed', 'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am',
          'Humidity3pm', 'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm',
          'Temp9am', 'Temp3pm', 'RISK_MM', 'WindGustDir', 'WindDir9am',
          'WindDir3pm', 'RainToday', 'RainTomorrow'],
         dtype='object')
```

```
[50]: # sns.pairplot(pairplot_last, diag_kind= 'kde')
```

```
[51]: processed_df_log
```

```
[51]:      MinTemp    MaxTemp    Rainfall    Evaporation    Sunshine    WindGustSpeed \
0      2.079442    3.190476         NaN        1.223775    1.840550        3.401197
1      2.639057    3.292126    1.280934        1.481605    2.272126        3.663562
2      2.617396    3.152736    1.280934        1.757858    1.193922        4.442651
3      2.587764    2.740840    3.683867        1.974081    2.208274        3.988984
4      2.028148    2.778819    1.029619        1.722767    2.360854        3.912023
..      ...      ...      ...      ...      ...      ...
361    2.197225    3.424263         NaN        2.028148    2.493205        4.330733
362    1.960095    3.346389         NaN        2.451005    2.541602        3.871201
363    2.525729    2.990720         NaN        2.128232    1.667707        3.761200
364    2.525729    3.292126         NaN        1.609438    1.960095        3.828641
365    2.509599    3.407842         NaN        1.791759    2.533697        4.356709
```

```
      WindSpeed9am    WindSpeed3pm    Humidity9am    Humidity3pm    ...    Cloud9am \
0          1.791759        2.995732        4.219508        3.367296    ...    1.945910
1          1.386294        2.833213        4.382027        3.583519    ...    1.609438
2          1.791759        1.791759        4.406719        4.234107    ...    2.079442
3          3.401197        3.178054        4.127134        4.025352    ...    0.693147
4          2.995732        3.332205        4.219508        3.891820    ...    1.945910
..      ...      ...      ...      ...      ...
361        1.945910        3.912023        3.637586        2.708050    ...    0.000000
362        0.693147        2.944439        3.806662        3.091042    ...         NaN
363        2.397895        2.197225        4.143135        3.850148    ...    1.098612
364        1.791759        3.332205        4.234107        3.663562    ...    1.791759
365        3.433987        3.555348        3.761200        2.564949    ...    0.000000
```

```
      Cloud3pm    Temp9am    Temp3pm    RISK_MM    WindGustDir    WindDir9am \
0      1.945910    2.667228    3.161247    1.280934        1.945910        2.484907
1      1.098612    2.862201    3.246491    1.280934        0.000000         NaN
2      1.945910    2.734368    3.005683    3.683867        1.945910        1.098612
3      1.945910    2.602690    2.646175    1.029619        1.945910        2.639057
4      1.945910    2.406945    2.734368         NaN        2.302585        2.302585
..      ...      ...      ...      ...      ...
361    1.098612    3.015535    3.401197         NaN        1.791759        2.302585
362    0.000000    2.844909    3.339322         NaN        1.098612        1.791759
363    0.693147    2.674149    2.906901         NaN        0.693147        0.000000
364    1.945910    2.760010    3.254243         NaN        1.945910        2.397895
365    0.000000    3.169686    3.353407         NaN        1.945910        1.945910
```

```
      WindDir3pm    RainToday    RainTomorrow
0          1.945910         NaN          0.0
1          2.564949          0.0          0.0
2          1.609438          0.0          0.0
3          2.564949          0.0          0.0
4          0.693147          0.0         NaN
..      ...      ...      ...
361        1.945910         NaN         NaN
```

```

362    1.791759      NaN      NaN
363    0.000000      NaN      NaN
364    2.639057      NaN      NaN
365    2.639057      NaN      NaN

```

[366 rows x 22 columns]

```

[52]: processed_df_log_new = processed_df_log.copy()
      for i in processed_df_log_new.columns:
          processed_df_log_new[i] = processed_df_log_new[i].replace(np.nan,
          ↪processed_df_log_new[i].mean())

```

```

[82]: # processed_df_log_new

```

```

[54]: W = processed_df_log_new.iloc[:, :-1]
      V = processed_df_log_new.iloc[:, -1]

```

```

[85]: # print("Data :\n", W, "\n")

```

```

[56]: W = whiten(W)

```

```

[57]: centroids, mean_dist = kmeans(W, 2)
      print("Code-book :\n", centroids, "\n")

```

Code-book :

```

[[2.52210223e+00 9.53420945e+00 1.01201283e+00 2.56172339e+00
 2.75176729e+00 1.05582082e+01 2.96989569e+00 5.13882326e+00
 2.17590122e+01 9.07017914e+00 1.05621531e+03 1.08825750e+03
 1.35172903e+00 1.38487484e+00 4.33229197e+00 8.53975666e+00
 1.06884954e+00 2.21303459e+00 2.95462803e+00 2.71024468e+00
 0.00000000e+00]
[1.33264334e+00 7.97570784e+00 7.86389251e-01 1.12030518e+00
 2.39969657e+00 9.90075228e+00 2.96679099e+00 5.14015655e+00
 2.24817321e+01 9.78932981e+00 1.05716151e+03 1.08928294e+03
 1.31231602e+00 1.32296791e+00 2.94173335e+00 7.03639851e+00
 7.2282725e-01 2.65915202e+00 3.41978484e+00 2.81204609e+00
 0.00000000e+00]]

```

```

[58]: clusters, dist = vq(W, centroids)
      print("Clusters :\n", clusters, "\n")

```

Clusters :

```

[0 0 0 0 0 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 1 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0

```



```

1 0 0 0 0 0 0 1 1 1 0 1 1 0 0 0 1 1 1 1 0 1 1 1 1 1 1 1 0 0 1 1 1 1 1 1
1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 0 0 1 1 1 0 0 0 0 0 1 1 0 0 0
0 0 0 0 0 0 0 0 1 1 1 0 0 0 0 0 0 1 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0]

```

```

[59]: cluster1 = list(clusters).count(0)

cluster2 = list(clusters).count(1)

```

```

[60]: kmeanss = KMeans(n_clusters=2, random_state=42)

```

```

[61]: silhouette_score(W, kmeanss.fit_predict(W))

```

```

[61]: 0.16255911532588155

```

```

[62]: colors = 10*["g","r","c","b","k"]

class K_Means:
    def __init__(self, k=2, tol=0.001, max_iter=300):
        self.k = k
        self.tol = tol
        self.max_iter = max_iter
    def fit(self,data):

        self.centroids = {}

        for i in range(self.k):
            self.centroids[i] = data[i]
        for i in range(self.max_iter):
            self.classifications = {}

            for i in range(self.k):
                self.classifications[i] = []

            for featureset in data:
                distances = [np.linalg.norm(featureset-self.
↪centroids[centroid]) for centroid in self.centroids]
                classification = distances.index(min(distances))
                self.classifications[classification].append(featureset)

            prev_centroids = dict(self.centroids)

            for classification in self.classifications:

```

```

        self.centroids[classification] = np.average(self.
↪classifications[classification],axis=0)

    optimized = True

    for c in self.centroids:
        original_centroid = prev_centroids[c]
        current_centroid = self.centroids[c]
        if np.sum((current_centroid-original_centroid)/
↪original_centroid*100.0) > self.tol:
            print(np.sum((current_centroid-original_centroid)/
↪original_centroid*100.0))
            optimized = False

    if optimized:
        break



    def predict(self,data):
        distances = [np.linalg.norm(data-self.centroids[centroid]) for
↪centroid in self.centroids]
        classification = distances.index(min(distances))
        return classification

clf = K_Means()
clf.fit(W)

for centroid in clf.centroids:
    plt.scatter(clf.centroids[centroid][0], clf.centroids[centroid][1],
        marker="o", color="k", s=150, linewidths=5)

for classification in clf.classifications:
    color = colors[classification]
    for featureset in clf.classifications[classification]:
        plt.scatter(featureset[0], featureset[1], marker="x", color=color,
↪s=150, linewidths=5)
    plt.show()

```



day-5

February 27, 2024

```
[5]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import load_wine
import warnings
warnings.filterwarnings('ignore')
```

```
[6]: # Load the data
data = load_wine()
```

```
[7]: data.keys()
```

```
[7]: dict_keys(['data', 'target', 'frame', 'target_names', 'DESCR', 'feature_names'])
```

```
[8]: # Description of the data
print(data.DESCR)
```

```
.. _wine_dataset:
```

```
Wine recognition dataset
```

```
-----
```

```
**Data Set Characteristics:**
```

```
:Number of Instances: 178
```

```
:Number of Attributes: 13 numeric, predictive attributes and the class
```

```
:Attribute Information:
```

- Alcohol
- Malic acid
- Ash
- Alcalinity of ash
- Magnesium
- Total phenols
- Flavanoids
- Nonflavanoid phenols
- Proanthocyanins

- Color intensity
- Hue
- OD280/OD315 of diluted wines
- Proline
- class:
 - class_0
 - class_1
 - class_2

:Summary Statistics:

	Min	Max	Mean	SD
Alcohol:	11.0	14.8	13.0	0.8
Malic Acid:	0.74	5.80	2.34	1.12
Ash:	1.36	3.23	2.36	0.27
Alcalinity of Ash:	10.6	30.0	19.5	3.3
Magnesium:	70.0	162.0	99.7	14.3
Total Phenols:	0.98	3.88	2.29	0.63
Flavanoids:	0.34	5.08	2.03	1.00
Nonflavanoid Phenols:	0.13	0.66	0.36	0.12
Proanthocyanins:	0.41	3.58	1.59	0.57
Colour Intensity:	1.3	13.0	5.1	2.3
Hue:	0.48	1.71	0.96	0.23
OD280/OD315 of diluted wines:	1.27	4.00	2.61	0.71
Proline:	278	1680	746	315

:Missing Attribute Values: None

:Class Distribution: class_0 (59), class_1 (71), class_2 (48)

:Creator: R.A. Fisher

:Donor: Michael Marshall (MARSHALL%PLU@io.arc.nasa.gov)

:Date: July, 1988

This is a copy of UCI ML Wine recognition datasets.

<https://archive.ics.uci.edu/ml/machine-learning-databases/wine/wine.data>

The data is the results of a chemical analysis of wines grown in the same region in Italy by three different cultivators. There are thirteen different measurements taken for different constituents found in the three types of wine.

Original Owners:

Forina, M. et al, PARVUS -

An Extendible Package for Data Exploration, Classification and Correlation.
Institute of Pharmaceutical and Food Analysis and Technologies,

Via Brigata Salerno, 16147 Genoa, Italy.

Citation:

Lichman, M. (2013). UCI Machine Learning Repository
[<https://archive.ics.uci.edu/ml>]. Irvine, CA: University of California,
School of Information and Computer Science.

|details-start|

****References****

|details-split|

(1) S. Aeberhard, D. Coomans and O. de Vel,
Comparison of Classifiers in High Dimensional Settings,
Tech. Rep. no. 92-02, (1992), Dept. of Computer Science and Dept. of
Mathematics and Statistics, James Cook University of North Queensland.
(Also submitted to Technometrics).

The data was used with many others for comparing various
classifiers. The classes are separable, though only RDA
has achieved 100% correct classification.
(RDA : 100%, QDA 99.4%, LDA 98.9%, 1NN 96.1% (z-transformed data))
(All results using the leave-one-out technique)

(2) S. Aeberhard, D. Coomans and O. de Vel,
"THE CLASSIFICATION PERFORMANCE OF RDA"
Tech. Rep. no. 92-01, (1992), Dept. of Computer Science and Dept. of
Mathematics and Statistics, James Cook University of North Queensland.
(Also submitted to Journal of Chemometrics).

|details-end|

```
[9]: # Features
      data.feature_names
```

```
[9]: ['alcohol',
      'malic_acid',
      'ash',
      'alcalinity_of_ash',
      'magnesium',
      'total_phenols',
      'flavanoids',
      'nonflavanoid_phenols',
      'proanthocyanins',
      'color_intensity',
      'hue',
```

```
'od280/od315_of_diluted_wines',
'proline']
```

```
[10]: # Create a dataframe with the data
df = pd.DataFrame(data.data, columns=data.feature_names)
df.head()
```

```
[10]:   alcohol  malic_acid  ash  alcalinity_of_ash  magnesium  total_phenols  \
0    14.23      1.71  2.43          15.6      127.0          2.80
1    13.20      1.78  2.14          11.2      100.0          2.65
2    13.16      2.36  2.67          18.6      101.0          2.80
3    14.37      1.95  2.50          16.8      113.0          3.85
4    13.24      2.59  2.87          21.0      118.0          2.80

      flavanoids  nonflavanoid_phenols  proanthocyanins  color_intensity  hue  \
0         3.06              0.28          2.29          5.64  1.04
1         2.76              0.26          1.28          4.38  1.05
2         3.24              0.30          2.81          5.68  1.03
3         3.49              0.24          2.18          7.80  0.86
4         2.69              0.39          1.82          4.32  1.04

      od280/od315_of_diluted_wines  proline
0              3.92      1065.0
1              3.40      1050.0
2              3.17      1185.0
3              3.45      1480.0
4              2.93       735.0
```

```
[11]: # Add the target and target names to the dataframe
df['target'] = data.target
df['target_names'] = df.target.apply(lambda x: data.target_names[x])
```

```
[12]: df.head().T
```

```
[12]:           0         1         2         3         4
alcohol    14.23    13.2    13.16    14.37    13.24
malic_acid   1.71     1.78     2.36     1.95     2.59
ash         2.43     2.14     2.67     2.5     2.87
alcalinity_of_ash  15.6    11.2    18.6    16.8    21.0
magnesium   127.0   100.0   101.0   113.0   118.0
total_phenols  2.8     2.65     2.8     3.85     2.8
flavanoids   3.06     2.76     3.24     3.49     2.69
nonflavanoid_phenols  0.28    0.26     0.3     0.24    0.39
proanthocyanins  2.29    1.28     2.81     2.18     1.82
color_intensity  5.64    4.38     5.68     7.8     4.32
hue          1.04    1.05     1.03     0.86     1.04
od280/od315_of_diluted_wines  3.92     3.4     3.17     3.45     2.93
```

proline	1065.0	1050.0	1185.0	1480.0	735.0
target	0	0	0	0	0
target_names	class_0	class_0	class_0	class_0	class_0

```
[13]: # Check for null values
df.isnull().sum()
```

```
[13]: alcohol          0
      malic_acid      0
      ash             0
      alcalinity_of_ash 0
      magnesium       0
      total_phenols   0
      flavanoids       0
      nonflavanoid_phenols 0
      proanthocyanins  0
      color_intensity 0
      hue              0
      od280/od315_of_diluted_wines 0
      proline          0
      target           0
      target_names     0
      dtype: int64
```

```
[14]: # Check the distribution of the target
df.target_names.value_counts()
```

```
[14]: target_names
      class_1    71
      class_0    59
      class_2    48
      Name: count, dtype: int64
```

```
[15]: # Check the distribution of the features
df.describe().T
```

	count	mean	std	min	25% \
alcohol	178.0	13.000618	0.811827	11.03	12.3625
malic_acid	178.0	2.336348	1.117146	0.74	1.6025
ash	178.0	2.366517	0.274344	1.36	2.2100
alcalinity_of_ash	178.0	19.494944	3.339564	10.60	17.2000
magnesium	178.0	99.741573	14.282484	70.00	88.0000
total_phenols	178.0	2.295112	0.625851	0.98	1.7425
flavanoids	178.0	2.029270	0.998859	0.34	1.2050
nonflavanoid_phenols	178.0	0.361854	0.124453	0.13	0.2700
proanthocyanins	178.0	1.590899	0.572359	0.41	1.2500
color_intensity	178.0	5.058090	2.318286	1.28	3.2200

hue	178.0	0.957449	0.228572	0.48	0.7825
od280/od315_of_diluted_wines	178.0	2.611685	0.709990	1.27	1.9375
proline	178.0	746.893258	314.907474	278.00	500.5000
target	178.0	0.938202	0.775035	0.00	0.0000

	50%	75%	max
alcohol	13.050	13.6775	14.83
malic_acid	1.865	3.0825	5.80
ash	2.360	2.5575	3.23
alcalinity_of_ash	19.500	21.5000	30.00
magnesium	98.000	107.0000	162.00
total_phenols	2.355	2.8000	3.88
flavanoids	2.135	2.8750	5.08
nonflavanoid_phenols	0.340	0.4375	0.66
proanthocyanins	1.555	1.9500	3.58
color_intensity	4.690	6.2000	13.00
hue	0.965	1.1200	1.71
od280/od315_of_diluted_wines	2.780	3.1700	4.00
proline	673.500	985.0000	1680.00
target	1.000	2.0000	2.00

```
[16]: df.dtypes
```

```
[16]: alcohol          float64
malic_acid          float64
ash                 float64
alcalinity_of_ash   float64
magnesium           float64
total_phenols       float64
flavanoids          float64
nonflavanoid_phenols float64
proanthocyanins     float64
color_intensity     float64
hue                 float64
od280/od315_of_diluted_wines float64
proline             float64
target              int64
target_names        object
dtype: object
```

```
[17]: obj_df = df.select_dtypes(include=['object']).copy()
num_df = df.select_dtypes(include=['float64', 'int64']).copy()
```

```
[18]: #label encoding
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
label_df = obj_df.apply(le.fit_transform)
```


```
[19]: label_df.head()
```

```
[19]: target_names
0      0
1      0
2      0
3      0
4      0
```

```
[20]: #concatenate the label and numerical dataframes
new_df = pd.concat([num_df, label_df], axis=1)
```

```
[21]: #findin optimal number of clusters
from sklearn.cluster import KMeans
wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i)
    kmeans.fit(new_df)
    wcss.append(kmeans.inertia_)
plt.plot(range(1, 11), wcss)
```


```
[21]: [<matplotlib.lines.Line2D at 0x76a4d016a250>]
```




```
[22]: #fitting kmeans to the dataset
kmeans = KMeans(n_clusters=3)
y_kmeans = kmeans.fit_predict(new_df)
```


```
[23]: plt.scatter(new_df.iloc[y_kmeans == 0, 0], new_df.iloc[y_kmeans == 0, 1], s = 100, c = 'red', label = 'Cluster 1')
plt.scatter(new_df.iloc[y_kmeans == 1, 0], new_df.iloc[y_kmeans == 1, 1], s = 100, c = 'blue', label = 'Cluster 2')
plt.scatter(new_df.iloc[y_kmeans == 2, 0], new_df.iloc[y_kmeans == 2, 1], s = 100, c = 'green', label = 'Cluster 3')
plt.scatter(kmeans.cluster_centers[:, 0], kmeans.cluster_centers[:, 1], s = 200, c = 'yellow', label = 'Centroids')
plt.legend()
```

```
[23]: <matplotlib.legend.Legend at 0x76a4cdb67ed0>
```



```
[24]: sns.scatterplot(x='alcohol', y='malic_acid', hue='target_names', data=df)
```


```
[24]: <Axes: xlabel='alcohol', ylabel='malic_acid'>
```



```
[25]: #silhouette score
from sklearn.metrics import silhouette_score
silhouette_score(new_df, y_kmeans)
```

[25]: 0.55956126417174


```
[26]: #hierarchical clustering
import scipy.cluster.hierarchy as sch
dendrogram = sch.dendrogram(sch.linkage(new_df, method='single'))
plt.title('Dendrogram')
plt.xlabel('Customers')
plt.ylabel('Euclidean distances')
plt.show()
```



```
[27]: #DBSCAN
from sklearn.cluster import DBSCAN
dbscan = DBSCAN(eps=3, min_samples=4)
y_dbscan = dbscan.fit_predict(new_df)
```

```
[28]: sns.scatterplot(x='alcohol', y='malic_acid', hue='target_names', data=df)
```

```
[28]: <Axes: xlabel='alcohol', ylabel='malic_acid'>
```



```
[29]: # #silhouette score  
# from sklearn.metrics import silhouette_score  
# silhouette_score(new_df, y_dbscan)
```

```
[ ]:
```

day-6

February 27, 2024

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: df = pd.read_csv('placement.csv')
```

```
[3]: df.isna().sum()
```

```
[3]: cgpa      0
package      0
dtype: int64
```

```
[4]: df.dtypes
```

```
[4]: cgpa      float64
package  float64
dtype: object
```

```
[5]: def outliers(df, ft):
    q1 = df[ft].quantile(0.25)
    q3 = df[ft].quantile(0.75)
    iqr = q3 - q1
    ub = q3 + 1.5 * iqr
    lb = q1 - 1.5 * iqr
    ls = df.index[(df[ft] > ub) | (df[ft] < lb)]
    return ls

def replace_outliers_with_mean(df, ft):
    q1 = df[ft].quantile(0.25)
    q3 = df[ft].quantile(0.75)
    iqr = q3 - q1
    ub = q3 + 1.5 * iqr
    lb = q1 - 1.5 * iqr
```

```

    outliers_indices = df.index[(df[ft] >= ub) | (df[ft] <= lb)]

    df.loc[outliers_indices, ft] = df[ft].mean()

    return df

index_ls = []

for i in df.columns:
    index_ls.extend(outliers(df, i))
print(index_ls)

for i in df.columns:
    df = replace_outliers_with_mean(df, i)

```

[]

```

[6]: # #standard scaling
      # from sklearn.preprocessing import StandardScaler
      # scaler = StandardScaler()
      # df = scaler.fit_transform(df)

```

```

[7]: #linear regression
      from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LinearRegression

      X = df.iloc[:, :-1]
      y = df.iloc[:, -1]

      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=42)

```

```

[8]: regressor = LinearRegression()
      regressor.fit(X_train, y_train)

```

```

[8]: LinearRegression()

```

```

[9]: y_pred = regressor.predict(X_test)

```

```


[10]: #r2 score
       from sklearn.metrics import r2_score
       r2 = r2_score(y_test, y_pred)
       print(r2)

```

0.7730984312051673


```
[11]: #boxplot
sns.boxplot(data=df)
```

```
[11]: <Axes: >
```




```
[12]: #distplot
sns.histplot(df)
```

```
[12]: <Axes: ylabel='Count'>
```



```
[13]: sns.distplot(df)
```


```
[13]: <Axes: ylabel='Density'>
```

```
[14]: #yeo-johnson transformation
from sklearn.preprocessing import PowerTransformer
pt = PowerTransformer()
new_df = pd.DataFrame(pt.fit_transform(df))
```

```
[15]: #boxplot
sns.boxplot(data=new_df)
```

```
[15]: <Axes: >
```



```
[16]: W = new_df.iloc[:, :-1]
      z = new_df.iloc[:, -1]

      W_train, W_test, z_train, z_test = train_test_split(W, z, test_size=0.2,
      ↪random_state=0)
```

```
[17]: regressor = LinearRegression()
      regressor.fit(W_train, z_train)

      z_pred = regressor.predict(W_test)

      #r2 score
      r2 = r2_score(z_test, z_pred)
      print(r2)
```

0.7307234904168096

```
[18]: #random forest regression
      from sklearn.ensemble import RandomForestRegressor
      regressor = RandomForestRegressor(n_estimators=10, random_state=42)
      regressor.fit(W_train, z_train)
```

```
[18]: RandomForestRegressor(n_estimators=10, random_state=42)
```


```
[19]: z_pred = regressor.predict(W_test)
```

```
#r2 score  
r2 = r2_score(z_test, z_pred)  
print(r2)
```

0.4254801167712502


```
[20]: #scatter plot with slope  
plt.scatter(z_test, z_pred)  
plt.plot(z_test, z_test, color='red')
```

```
[20]: [<matplotlib.lines.Line2D at 0x77b65363f790>]
```



```
[21]: #scatter plot with slope  
plt.scatter(y_test, y_pred)  
plt.plot(y_test, y_test, color='red')  
print('Coefficient of determination: %.2f' % r2_score(y_test, y_pred))
```

Coefficient of determination: 0.77



[]:

day-7

February 27, 2024

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: df = pd.read_csv('auto-mpg.csv')
```

```
[3]: df.head()
```

```
[3]:      mpg  cylinders  displacement  horsepower  weight  acceleration  model year  \
0   18.0         8         307.0         130    3504         12.0         70
1   15.0         8         350.0         165    3693         11.5         70
2   18.0         8         318.0         150    3436         11.0         70
3   16.0         8         304.0         150    3433         12.0         70
4   17.0         8         302.0         140    3449         10.5         70
```

```
      origin  car name
0         1  chevrolet chevelle malibu
1         1      buick skylark 320
2         1  plymouth satellite
3         1      amc rebel sst
4         1      ford torino
```

```
[4]: df.dtypes
```

```
[4]: mpg          float64
cylinders        int64
displacement     float64
horsepower       object
weight           int64
acceleration     float64
model year       int64
origin           int64
car name         object
dtype: object
```

```
[5]: obj_df = df.select_dtypes(include=['object']).copy()
num_df = df.select_dtypes(include=['float64', 'int64']).copy()
```

```
[6]: obj_df.head().T
```

```
[6]:
           0           1           2 \
horsepower          130          165          150
car name  chevrolet chevelle malibu  buick skylark 320  plymouth satellite

           3           4
horsepower          150          140
car name  amc rebel sst  ford torino
```

```
[7]: num_df.head().T
```

```
[7]:
           0           1           2           3           4
mpg          18.0          15.0          18.0          16.0          17.0
cylinders          8.0           8.0           8.0           8.0           8.0
displacement  307.0          350.0          318.0          304.0          302.0
weight       3504.0          3693.0          3436.0          3433.0          3449.0
acceleration   12.0           11.5           11.0           12.0           10.5
model year      70.0           70.0           70.0           70.0           70.0
origin          1.0           1.0           1.0           1.0           1.0
```

```
[8]: #IQR method
Q1 = num_df.quantile(0.25)
Q3 = num_df.quantile(0.75)
IQR = Q3 - Q1

print(IQR)
```

```
mpg          11.50
cylinders          4.00
displacement   157.75
weight       1384.25
acceleration    3.35
model year        6.00
origin           1.00
dtype: float64
```

```
[9]: print((num_df < (Q1 - 1.5 * IQR)) | (num_df > (Q3 + 1.5 * IQR)))
```

```

    mpg  cylinders  displacement  weight  acceleration  model year  origin
0  False      False          False  False           False      False  False
1  False      False          False  False           False      False  False
2  False      False          False  False           False      False  False
3  False      False          False  False           False      False  False
4  False      False          False  False           False      False  False
```

```

..      ...      ...      ...      ...      ...      ...
393  False      False      False  False  False      False  False  False
394  False      False      False  False  False      True   False  False
395  False      False      False  False  False      False  False  False
396  False      False      False  False  False      False  False  False
397  False      False      False  False  False      False  False  False

```

[398 rows x 7 columns]

```

[10]: #Removing outliers
num_df_out = num_df[~((num_df < (Q1 - 1.5 * IQR)) | (num_df > (Q3 + 1.5 *
↪IQR))).any(axis=1)]

```

```

[11]: obj_df.head().T

```

```

[11]:
           0           1           2 \
horsepower          130          165          150
car name    chevrolet chevelle malibu  buick skylark 320  plymouth satellite

           3           4
horsepower          150          140
car name    amc rebel sst  ford torino

```

```

[12]: #label encoding
from sklearn.preprocessing import LabelEncoder

lb_make = LabelEncoder()
label_obj_df = obj_df.apply(lb_make.fit_transform)

```

```

[13]: label_obj_df.head().T

```

```

[13]:
           0    1    2    3    4
horsepower  15  33  27  27  22
car name    49  36  231  14  161

```

```

[14]: new_df = pd.concat([num_df_out, label_obj_df], axis=1)

```

```

[15]: new_df.head().T

```

```

[15]:
           0           1           2           3           4
mpg        18.0        15.0        18.0        16.0        17.0
cylinders    8.0         8.0         8.0         8.0         8.0
displacement 307.0      350.0      318.0      304.0      302.0
weight      3504.0     3693.0     3436.0     3433.0     3449.0
acceleration  12.0      11.5       11.0       12.0      10.5
model year   70.0      70.0       70.0       70.0      70.0
origin        1.0        1.0         1.0         1.0         1.0

```

horsepower	15.0	33.0	27.0	27.0	22.0
car name	49.0	36.0	231.0	14.0	161.0

```
[16]: new_df.isna().sum()
```

```
[16]: mpg            8
      cylinders      8
      displacement  8
      weight        8
      acceleration  8
      model year    8
      origin        8
      horsepower    0
      car name      0
      dtype: int64
```

```
[17]: mena_num_df = new_df.fillna(new_df.mean())
```

lasso

```
[18]: #lasso regression
      from sklearn.model_selection import train_test_split
      from sklearn.linear_model import Lasso
      from sklearn.metrics import mean_squared_error

      X = mena_num_df.drop('mpg', axis=1)
      y = mena_num_df['mpg']

      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
      ↪random_state=42)
```

```
[19]: lasso = Lasso(alpha=0.1)
      lasso.fit(X_train, y_train)
```

```
[19]: Lasso(alpha=0.1)
```

```
[20]: y_pred = lasso.predict(X_test)
```

```
[21]: print(mean_squared_error(y_test, y_pred))
```

```
10.834734885814798
```

```
[22]: #r2 score
      from sklearn.metrics import r2_score

      print(r2_score(y_test, y_pred))
```

```
0.8258432055490966
```


Ridge

```
[23]: #ridge regression
      from sklearn.linear_model import Ridge

      ridge = Ridge(alpha=1)
      ridge.fit(X_train, y_train)
```

```
[23]: Ridge(alpha=1)
```

```
[24]: y_pred = ridge.predict(X_test)
```

```
[25]: print(mean_squared_error(y_test, y_pred))
```

10.746788369431064

```
[26]: print(r2_score(y_test, y_pred))
```

0.8272568518946632

ElasticNet

```
[27]: # elastic net
      from sklearn.linear_model import ElasticNet

      elastic = ElasticNet(alpha=0.1, l1_ratio=0.5)
      elastic.fit(X_train, y_train)
```

```
[27]: ElasticNet(alpha=0.1)
```

```
[28]: y_pred = elastic.predict(X_test)
```

```
[29]: print(mean_squared_error(y_test, y_pred))
```

10.830856032105528

```
[30]: print(r2_score(y_test, y_pred))
```

0.825905553980809

```
[ ]:
```

day7-2

February 27, 2024

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import PowerTransformer
import warnings
warnings.filterwarnings('ignore')
```

```
[3]: df = pd.read_csv('train .csv')
```

```
[4]: df.head().T
```

```
[4]:
```

	0	1	2 \
employee_id	65438	65141	7513
department	Sales & Marketing	Operations	Sales & Marketing
region	region_7	region_22	region_19
education	Master's & above	Bachelor's	Bachelor's
gender	f	m	m
recruitment_channel	sourcing	other	sourcing
no_of_trainings	1	1	1
age	35	30	34
previous_year_rating	5.0	5.0	3.0
length_of_service	8	4	7
KPIs_met >80%	1	0	0
awards_won?	0	0	0
avg_training_score	49	60	50
is_promoted	0	0	0

	3	4
employee_id	2542	48945
department	Sales & Marketing	Technology
region	region_23	region_26
education	Bachelor's	Bachelor's
gender	m	m

recruitment_channel	other	other
no_of_trainings	2	1
age	39	45
previous_year_rating	1.0	3.0
length_of_service	10	2
KPIs_met >80%	0	0
awards_won?	0	0
avg_training_score	50	73
is_promoted	0	0

```
[5]: df.dtypes
```

```
[5]: employee_id      int64
      department      object
      region          object
      education        object
      gender          object
      recruitment_channel  object
      no_of_trainings  int64
      age             int64
      previous_year_rating float64
      length_of_service  int64
      KPIs_met >80%     int64
      awards_won?      int64
      avg_training_score int64
      is_promoted      int64
      dtype: object
```

```
[6]: obj_df = df.select_dtypes(include=['object']).copy()
      num_df = df.select_dtypes(include=['int64']).copy()
```

```
[7]: obj_df.head().T
```

```
[7]:
```

	0	1	2 \
department	Sales & Marketing	Operations	Sales & Marketing
region	region_7	region_22	region_19
education	Master's & above	Bachelor's	Bachelor's
gender	f	m	m
recruitment_channel	sourcing	other	sourcing

	3	4
department	Sales & Marketing	Technology
region	region_23	region_26
education	Bachelor's	Bachelor's
gender	m	m
recruitment_channel	other	other

```
[8]: num_df.head().T
```

```
[8]:
```

	0	1	2	3	4
employee_id	65438	65141	7513	2542	48945
no_of_trainings	1	1	1	2	1
age	35	30	34	39	45
length_of_service	8	4	7	10	2
KPIs_met >80%	1	0	0	0	0
awards_won?	0	0	0	0	0
avg_training_score	49	60	50	50	73
is_promoted	0	0	0	0	0

```
[9]: obj_df.isnull().sum()
```

```
[9]:
```

department	0
region	0
education	2409
gender	0
recruitment_channel	0

dtype: int64

```
[10]: num_df.isnull().sum()
```

```
[10]:
```

employee_id	0
no_of_trainings	0
age	0
length_of_service	0
KPIs_met >80%	0
awards_won?	0
avg_training_score	0
is_promoted	0

dtype: int64

```
[11]: #IQR
Q1 = num_df.quantile(0.25)
Q3 = num_df.quantile(0.75)

IQR = Q3 - Q1
print(IQR)
```

employee_id	39060.75
no_of_trainings	0.00
age	10.00
length_of_service	4.00
KPIs_met >80%	1.00
awards_won?	0.00
avg_training_score	25.00

```
is_promoted          0.00
dtype: float64
```

```
[12]: #Outliers
num_df_out = num_df[~((num_df < (Q1 - 1.5 * IQR)) | (num_df > (Q3 + 1.5 * IQR)))
      ↪any(axis=1)]
```

```
[13]: num_df_out.shape
```

```
[13]: (36477, 8)
```

```
[14]: #label encoding
from sklearn.preprocessing import LabelEncoder
labelencoder = LabelEncoder()
label_df = obj_df.apply(labelencoder.fit_transform)
```

```
[15]: label_df.head().T
```

```
[15]:
```

	0	1	2	3	4
department	7	4	7	7	8
region	31	14	10	15	18
education	2	0	0	0	0
gender	0	1	1	1	1
recruitment_channel	2	0	2	0	0

```
[16]: new_df = pd.concat([label_df, num_df], axis=1)
```

```
[17]: new_df.isna().sum()
```

```
[17]: department          0
region                 0
education              0
gender                 0
recruitment_channel    0
employee_id            0
no_of_trainings        0
age                    0
length_of_service      0
KPIs_met >80%          0
awards_won?            0
avg_training_score      0
is_promoted            0
dtype: int64
```

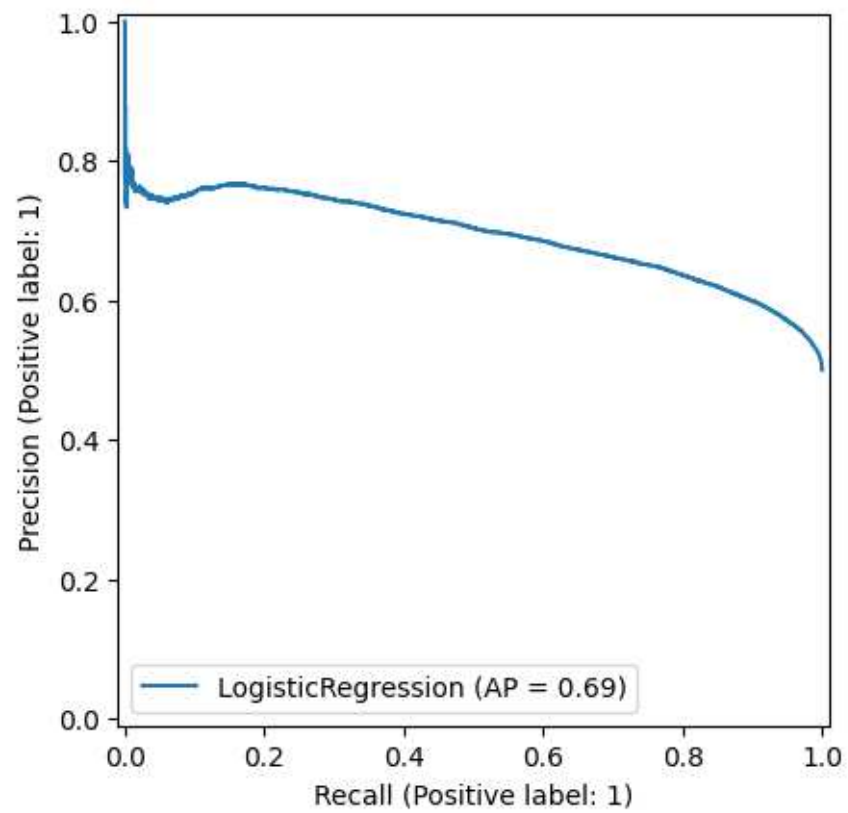
```
[18]: new_df.columns
```


```
[18]: Index(['department', 'region', 'education', 'gender', 'recruitment_channel',  
          'employee_id', 'no_of_trainings', 'age', 'length_of_service',  
          'KPIs_met >80%', 'awards_won?', 'avg_training_score', 'is_promoted'],  
         dtype='object')
```

```
[19]: X = new_df.drop(['is_promoted'], axis=1)  
y = new_df['is_promoted']  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
                                                    random_state=42)  
sm = SMOTE(random_state=42)  
X_train, y_train = sm.fit_resample(X_train, y_train)
```

```
[20]: #ROS  
from imblearn.over_sampling import RandomOverSampler  
from sklearn.metrics import PrecisionRecallDisplay  
from sklearn import linear_model  
  
oversampler = RandomOverSampler(sampling_strategy='minority')  
X_over, y_over = oversampler.fit_resample(X_train, y_train)  
  
lr = linear_model.LogisticRegression()  
lr.fit(X_over, y_over)  
y_pred = lr.predict(X_test)  
pr = PrecisionRecallDisplay.from_estimator(lr, X_over, y_over)  
pr.plot()
```

```
[20]: <sklearn.metrics._plot.precision_recall_curve.PrecisionRecallDisplay at  
0x76d9650cb410>
```






```
[21]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score, \
      ↪ roc_auc_score

print(mean_absolute_error(y_test, y_pred))
print(mean_squared_error(y_test, y_pred))
print(r2_score(y_test, y_pred))
print(roc_auc_score(y_test, y_pred))

#Draw ROC curve
from sklearn.metrics import roc_curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
plt.plot(fpr, tpr)
plt.show()
```

```
0.36717752234993617
0.36717752234993617
-3.833157557121062
0.5994591759564432
```

```
[22]: print(y_over.value_counts())
```

```
is_promoted
0    40086
1    40086
Name: count, dtype: int64
```

```
[23]: #find the best model
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix

knn = KNeighborsClassifier()
param_grid = {'n_neighbors': np.arange(1, 25)}
knn_gscv = GridSearchCV(knn, param_grid, cv=5)
knn_gscv.fit(X_over, y_over)
knn_gscv.best_params_


knn_gscv = KNeighborsClassifier(n_neighbors=24)
knn_gscv.fit(X_over, y_over)
y_pred = knn_gscv.predict(X_test)
print(confusion_matrix(y_test, y_pred))
```


```
print(classification_report(y_test, y_pred))
pr = PrecisionRecallDisplay.from_estimator(knn_gscv, X_over, y_over)
pr.plot()
```

```
[[6144 3910]
 [ 500  408]]
```

	precision	recall	f1-score	support
0	0.92	0.61	0.74	10054
1	0.09	0.45	0.16	908
accuracy			0.60	10962
macro avg	0.51	0.53	0.45	10962
weighted avg	0.86	0.60	0.69	10962

[23]: <sklearn.metrics._plot.precision_recall_curve.PrecisionRecallDisplay at 0x76d964b34cd0>





[]: